

Survival Analysis with Apache Spark and Apache SystemML on Stackexchange

Mateo Álvarez Calvo

May 13, 2017

Contents

1	Introduction & main goals	2
1.1	Main technologies	3
1.1.1	Apache Spark	3
1.1.2	Apache SystemML	3
1.2	The Stackexchange data	3
1.3	Main Objectives	4
2	Technologies	4
2.1	Apache Spark	4
2.1.1	Apache Spark Structure	4
2.2	Apache SystemML	4
2.2.1	SystemML Structure	5
2.3	Reproducible research with Python, Scala and Jupyter	5
3	Infrastructure and resources	5
3.1	Architecture scheme	5
3.2	Configuration	5
3.3	Workflow	5
4	Results	5
4.1	Used data	5
4.2	Kaplan-Meier model	5
4.3	Cox Proportional Hazards Model	5
5	Conclusions	5
5.1	Most important results	5
5.2	Lessons learnt	5

1 Introduction & main goals

Many studies have been done over the Stackexchange community [such as], one of the biggest Q&A sites in the world. The present is yet another study over the data of the famous site, but in this case, the study has two particularities, the use of Apache Spark with the library of Apache SystemML for the processing in a parallel environment, and the use of Survival Analysis to analyze the impact of the variables in the time an answer is accepted for each question, the "survival of each question" in the community.

1.1 Main technologies

As one of the biggest Q&A communities, Stackexchange has large amount of data of each interaction. Stackexchange is separated in several communities, regarding different topics. These communities can be small, as [] or really big, as Stackoverflow, the developers community. This particularity makes necessary the use technologies prepared to process large amounts of data, in the later case.

Among the available big data technologies, the

For this purpose, Apache Spark, the latest distributed open-source processing technology, has been chosen to parallelize the operations on the data.

Spark ML is the machine learning library of spark, which contains lots of algorithms. It also includes some Survival Analysis algorithms, but just for parametric modelling. This gives an excuse to use the recently adopted by the Apache Foundation SystemML, a machine learning library developed by IBM, which has non, semi and parametric algorithms for survival analysis.

1.1.1 Apache Spark

Apache Spark is a distributed processing technology developed in Scala by Databricks that represents the next step of Apache Hadoop, including the best parts of it, such as the Hadoop File System, but under a complete new paradigm that allows operations different from the famous map-reduce, using RAM as storage for results rather than writting to disk and lazy and optimized execution of tasks, to mention some of the main features.

1.1.2 Apache SystemML

Recently included in the Apache Foundation Incubating program, Apache SystemML is a machine learning library that works over distributed frameworks, Spark or Hadoop, written in a self made language, Distributed Machine Learning (dml) and with a python, scala or R like syntaxis API.

This library provides many distributed implementations of important algorithms as well as a syntax to create new algorithms in a distributed mode.

1.2 The Stackexchange data

Stackexchange facilitates raw partially processed data from the communities database every once in a while or every three months for data scientists and people in general to download and analyze. The data is available in a torrent file and each package has about 35 - 40 GB of compressed information.

This compressed file has data from different communities for a certain period of time. In our case, the analysis is done over Scyfy community, which is a median size community for science-fiction Q&A.

1.3 Main Objectives

For this study several objectives have been proposed:

- Use Spark to make the data cleaning to create a script for further research on the Stackexchange site.
- Verify SystemML integration with Spark for further research and scalability.
- Use SystemML survival analysis algorithms to analyze Stackexchange's data and obtain conclusions on the main variables affecting the time taken by the community to answer a question.

2 Technologies

The downloaded data from Stackexchange for the analysis [weights?] 40 GB, which is enough amount [size] to consider distributed processing.

2.1 Apache Spark

The latest distributed processing technology for data processing is Apache Spark, born in [2015] in [Ucla?], which is the predecessor of the well known processing framework Apache Hadoop, with all its ecosystem.

Comparing to Hadoop, Spark is able to speedup calculations using RAM instead of disk, with a higher variety of operations, not limited to Map-Reduce, and [].

2.1.1 Apache Spark Structure

2.2 Apache SystemML

Developed by IBM, Apache SystemML is the [precursor] of Spark ML (MLLib), the machine learning libraries of Spark. SystemML is coded in a self-made programming language, Distributed Machine Learning (DML) and it can be used from Spark or Hadoop in a [submit-like execution or in a interactive execution, the one it has been used in this study]

2.2.1 SystemML Structure

2.3 Reproducible research with Python, Scala and Jupyter

3 Infrastructure and resources

3.1 Architecture scheme

3.2 Configuration

3.3 Workflow

4 Results

4.1 Used data

4.2 Kaplan-Meier model

4.3 Cox Proportional Hazards Model

5 Conclusions

5.1 Most important results

5.2 Lessons learnt